

# Jamboree Case Study - Linear Regression

## Problem Statement

Jamboree recently launched a feature where students/learners can come to their website and check their probability of getting into the IVY league college. This feature estimates the chances of graduate admission from an Indian perspective.

### Additional View

- Lin Reg. will also help predict one's chances of admission given the rest of the variables.
- GRE Score, TOEFL Score & CGPA are most important attributes as per Indian Perspective.

## Installing Dependencies

In [709]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

## Loading Dataset

In [710]:

```
jamboree = pd.read_csv('https://d2beiqkhq929f0.cloudfront.net/public_assets/assets/000/0
```

In [711]:

```
jamboree.head(5)
```

Out[711]:

	Serial No.	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research	Chance of Admit
0	1	337	118	4	4.5	4.5	9.65	1	0.92
1	2	324	107	4	4.0	4.5	8.87	1	0.76
2	3	316	104	3	3.0	3.5	8.00	1	0.72
3	4	322	110	3	3.5	2.5	8.67	1	0.80
4	5	314	103	2	2.0	3.0	8.21	0	0.65

In [712]:

```
jamboree.drop(["Serial No."], axis = 1, inplace = True)
```

In [713]:

```
jamboree.shape
```

Out[713]:

```
(500, 8)
```

In [714]:

```
jamboree.dtypes
```

Out[714]:

```
GRE Score      int64
TOEFL Score    int64
University Rating  int64
SOP            float64
LOR            float64
CGPA           float64
Research        int64
Chance of Admit  float64
dtype: object
```

- All the features are numerical

In [715]:

```
jamboree.isnull().sum()
```

Out[715]:

```
GRE Score      0
TOEFL Score    0
University Rating  0
SOP            0
LOR            0
CGPA           0
Research        0
Chance of Admit  0
dtype: int64
```

- There is no null values

In [716]:

```
jamboree.describe().T
```

Out[716]:

	count	mean	std	min	25%	50%	75%	max
<b>GRE Score</b>	500.0	316.47200	11.295148	290.00	308.0000	317.00	325.00	340.00
<b>TOEFL Score</b>	500.0	107.19200	6.081868	92.00	103.0000	107.00	112.00	120.00
<b>University Rating</b>	500.0	3.11400	1.143512	1.00	2.0000	3.00	4.00	5.00
<b>SOP</b>	500.0	3.37400	0.991004	1.00	2.5000	3.50	4.00	5.00
<b>LOR</b>	500.0	3.48400	0.925450	1.00	3.0000	3.50	4.00	5.00
<b>CGPA</b>	500.0	8.57644	0.604813	6.80	8.1275	8.56	9.04	9.92
<b>Research</b>	500.0	0.56000	0.496884	0.00	0.0000	1.00	1.00	1.00
<b>Chance of Admit</b>	500.0	0.72174	0.141140	0.34	0.6300	0.72	0.82	0.97

- While Observing the mean and 50% percentile of data there is no significant difference observed
- We can conclude there are no outliers in the dataset.

In [717]:

```
jamboree.duplicated().sum()
```

Out[717]:

0

- There is no duplicated values in the dataset

In [718]:

```
jamboree.columns = map(lambda x: x.strip(), jamboree.columns)
```

## Non-Graphical Analysis

In [719]:

```
jamboree["University Rating"].value_counts(normalize=True)
```

Out[719]:

```
3    0.324
2    0.252
4    0.210
5    0.146
1    0.068
Name: University Rating, dtype: float64
```

- While observing the university rating. Most of universities average rated.

In [720]:

```
jamboree["SOP"].value_counts(normalize=True)
```

Out[720]:

```
4.0    0.178
3.5    0.176
3.0    0.160
2.5    0.128
4.5    0.126
2.0    0.086
5.0    0.084
1.5    0.050
1.0    0.012
```

Name: SOP, dtype: float64

In [721]:

```
jamboree["Research"].value_counts(normalize=True)
```

Out[721]:

```
1    0.56
0    0.44
```

Name: Research, dtype: float64

- Above stats shows there are almost equal distribution among students who did research

## Graphical Analysis: Univariate

In [722]:

```

num_cat = ["GRE Score", "TOEFL Score", "University Rating", "SOP", "LOR", "CGPA", "Resea





for i in range(len(num_cat)):
    fig = plt.figure(figsize = (15, 8))







    ax1 = plt.subplot2grid((1, 2), (0, 0))
    ax1.set_title(f"Distribution of {num_cat[i]}")
    ax1.set_xlabel(ax1.get_xlabel(), fontsize = 15)
    ax1.axvline(jamboree[num_cat[i]].mean(),color="green", label = "mean")
    ax1.axvline(jamboree[num_cat[i]].median(),color="red", label = "median")
    ax1.legend(loc = "best")
    sns.histplot(data=jamboree, x=num_cat[i], ax=ax1, kde=True)







    ax2 = plt.subplot2grid((1, 2), (0, 1))
    ax2.set_title(f"Box plot for {num_cat[i]}")
    ax2.set_xlabel(ax1.get_xlabel(), fontsize = 15)
    sns.boxplot(data = jamboree, x=num_cat[i], ax=ax2)

plt.show()

```







### • GRE Score Analysis

- Distribution of GRE resembles like Gaussian
- Mean of GRE Score is approx 315

- There is no outliers detected as mean and median overlaps
- **TOEFL Score Analysis**
  - Distribution of TOEFL somewhat resembles like Gaussian
  - Mean of TOEFL Score is approx 108
  - There is no outliers detected as mean and median overlaps
- **LOR Analysis**
  - Most of the students gets 3.5 out 5
- **CGPA Analysis**
  - Distribution of CGPA resembles like Gaussian
  - Mean of CGPA Score is approx 8.5
  - There is no outliers detected as mean and median overlaps
- **Chance of Admit**
  - Mean of chance of admission is 0.72

## Outliers Detection

- **From the above observation, There is no outliers detected in the dataset**

In [ ]:

## Graphical Analysis: Bivariate



In [723]:

```

num_cat = ["GRE Score", "TOEFL Score", "University Rating", "SOP", "LOR", "CGPA"]
target_cat = "Chance of Admit"





for i in range(len(num_cat)):
    fig = plt.figure(figsize = (15, 8))







    ax1 = plt.subplot2grid((1, 2), (0, 0))
    ax1.set_title(f"Distribution of {num_cat[i]}")
    ax1.set_xlabel(ax1.get_xlabel(), fontsize = 15)
    # sns.histplot(data=jamboree, x=num_cat[i], ax=ax1, kde=True)
    sns.regplot(data=jamboree, ax=ax1, x=num_cat[i], y=target_cat)

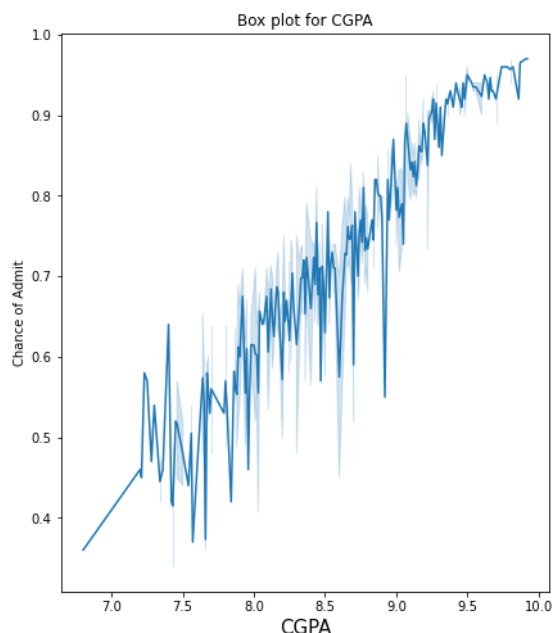

    ax2 = plt.subplot2grid((1, 2), (0, 1))
    ax2.set_title(f"Box plot for {num_cat[i]}")
    ax2.set_xlabel(ax1.get_xlabel(), fontsize = 15)
    # sns.boxplot(data = jamboree, x=num_cat[i], ax=ax2)
    sns.lineplot(data=jamboree, x=num_cat[i], y=target_cat, ax=ax2)

plt.show()

```







- **GRE vs Chance of Admit Analysis**

- There is linear relationship between GRE and Chance of Admission
- Higher the GRE -> Higher the chance of admission

- **TOEFL vs Chance of Admit Analysis**

- There is linear relationship between TOEFL and Chance of Admission
- Higher the TOEFL -> Higher the chance of admission

- **LOR / SOP / University Rating vs Chance of Admit Analysis**

- There is no significant linear relationship between TOEFL and Chance of Admission



- **CGPA vs Chance of Admit Analysis**

- There is linear relationship between TOEFL and Chance of Admission
- Higher the CGPA -> Higher the chance of admission

In [724]:

```
fig = sns.regplot(x="GRE Score",y="TOEFL Score",data=jamboree)
plt.title("GRE Score vs TOEFL Score")
plt.show()


fig = sns.regplot(x="GRE Score",y="CGPA",data=jamboree)
plt.title("GRE Score vs CGPA")
plt.show()
```



- People with higher GRE Scores also have higher TOEFL Scores which is justified because both TOEFL and GRE have a verbal section which although not similar are relatable
- Although there are exceptions, people with higher CGPA usually have higher GRE scores maybe because they are smart or hard working

In [725]:


```
fig = sns.scatterplot(x="CGPA", y="LOR", data=jamboree, hue="Research")
plt.title("LOR vs CGPA")
plt.show()
```



- LORs are not that related with CGPA so it is clear that a persons LOR is not dependent on that persons academic excellence.
- Having research experience is usually related with a good LOR which might be justified by the fact that supervisors have personal interaction with the students performing research which usually results in good LORs

In [726]:

```
fig = sns.scatterplot(x="GRE Score", y="LOR", data=jamboree, hue="Research")
plt.title("GRE Score vs CGPA")
plt.show()
```





- GRE scores and LORs are also not that related. People with different kinds of LORs have all kinds of GRE scores


In [727]:

```
fig = sns.scatterplot(x="CGPA", y="SOP", data=jamboree)
plt.title("SOP vs CGPA")
plt.show()

fig = sns.scatterplot(x="GRE Score", y="SOP", data=jamboree)
plt.title("GRE Score vs SOP")
plt.show()

fig = sns.scatterplot(x="TOEFL Score", y="SOP", data=jamboree)
plt.title("TOEFL Score vs SOP")
plt.show()
```






nt of Purpose is related to academic  
to be more hard working so they have good  
ght move towards higher CGPA as along

kinds of TOEFL Score. So the quality of  
kills.

## Correlation Analysis: Heat Map

In [728]:

```
fig = plt.figure(figsize = (15, 8))
corr = jamboree.corr()
sns.heatmap(corr, linewidths=.5, annot=True, cmap="Blues")
plt.show()
```



### • High Correlation

1. GRE Score vs TOEFL Score
2. CGPA vs TOEFL Score
3. CGPA vs GRE Score
4. Chance of Admit vs CGPA
5. GRE Score vs Chance of Admit

## Data Preprocessing

## Duplicate Value Check

In [729]:

```
np.any(jamboree.duplicated())
```

Out[729]:

False

## Missing Value Check

In [730]:

```
jamboree.isnull().sum()
```

Out[730]:

GRE Score	0
TOEFL Score	0
University Rating	0
SOP	0
LOR	0
CGPA	0
Research	0
Chance of Admit	0

dtype: int64

## Outlier Check



In [731]:

```
for i in data.columns:
    print("=====" * 10)
    print("Mean of {}: ".format(i), jamboree[i].mean())
    print("Median of {}: ".format(i), jamboree[i].median())
```

```
=====
Mean of GRE Score:  316.472
Median of GRE Score:  317.0
=====
Mean of TOEFL Score:  107.192
Median of TOEFL Score:  107.0
=====
Mean of University Rating:  3.114
Median of University Rating:  3.0
=====
Mean of SOP:  3.374
Median of SOP:  3.5
=====
Mean of LOR:  3.484
Median of LOR:  3.5
=====
Mean of CGPA:  8.576440000000003
Median of CGPA:  8.56
=====
Mean of Research:  0.56
Median of Research:  1.0
=====
Mean of Chance of Admit:  0.7217399999999996
Median of Chance of Admit:  0.72
```

- **No outliers detected. As each and every feature overlaps its mean and median**

## Feature Engineering & Data Modelling

In [732]:

```
from sklearn.model_selection import train_test_split

X = jamboree.drop(['Chance of Admit'], axis=1)
y = jamboree['Chance of Admit']

print("X shape: {}".format(X.shape))
print("y shape: {}".format(y.shape))
```

```
X shape: (500, 7)
y shape: (500,)
```

## Train & Test Split

In [733]:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, shuffle=True)

print("X_train shape: {}".format(X_train.shape))
print("X_test shape: {}".format(X_test.shape))
print("y_train shape: {}".format(y_train.shape))
print("y_test shape: {}".format(y_test.shape))
```

```
X_train shape: (400, 7)
X_test shape: (100, 7)
y_train shape: (400,)
y_test shape: (100,)
```

In [734]:

```
X_train.head(5)
```

Out[734]:

	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research
312	311	107	4	4.5	4.5	9.00	1
32	338	118	4	3.0	4.5	9.40	1
194	316	109	3	3.5	3.0	8.76	0
490	307	105	2	2.5	4.5	8.12	1
470	320	110	5	4.0	4.0	9.27	1

In [735]:

```
y_train
```

Out[735]:

```
312    0.78
32     0.91
194    0.77
490    0.67
470    0.87
...
145    0.81
492    0.53
59     0.42
339    0.81
88     0.64
Name: Chance of Admit, Length: 400, dtype: float64
```

## Feature standardization

In [736]:

```
# Standarization
from sklearn.preprocessing import StandardScaler
X_train_columns=X_train.columns
std=StandardScaler()
X_train_std=std.fit_transform(X_train)
```

In [737]:

```
X_train_std
```

Out[737]:

```
array([[ -0.48601649, -0.0276501 ,  0.79650712, ...,  1.11439835,
         0.695962 ,  0.86413245],
       [  1.87625988,  1.761474 ,  0.79650712, ...,  1.11439835,
         1.36072183,  0.86413245],
       [ -0.0485579 ,  0.29764519, -0.08361124, ..., -0.53047007,
         0.2971061 , -1.15723001],
       ...,
       [ -0.48601649, -0.51559304, -0.96372961, ..., -1.62704901,
        -0.46736771, -1.15723001],
       [  0.65137584, -0.0276501 ,  1.67662549, ...,  0.56610888,
         0.13091614,  0.86413245],
       [ -0.22354134,  0.13499755, -0.08361124, ...,  0.01781941,
        -0.73327164, -1.15723001]])
```

In [738]:

```
X_train = pd.DataFrame(data = X_train_std, columns = X_train_columns)
X_train.head(5)
```

Out[738]:

	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research
0	-0.486016	-0.027650	0.796507	1.119664	1.114398	0.695962	0.864132
1	1.876260	1.761474	0.796507	-0.379886	1.114398	1.360722	0.864132
2	-0.048558	0.297645	-0.083611	0.119964	-0.530470	0.297106	-1.157230
3	-0.835983	-0.352945	-0.963730	-0.879736	1.114398	-0.766510	0.864132
4	0.301409	0.460293	1.676625	0.619814	0.566109	1.144675	0.864132

## Model Building: Lin Reg, Lasso Reg, Ridge Reg

In [739]:

```
from sklearn.linear_model import LinearRegression, Lasso, Ridge
from sklearn.metrics import mean_squared_error
from statsmodels.stats.outliers_influence import variance_inflation_factor
import statsmodels.api as sm

lm = LinearRegression()
lm.fit(X_train.values, y_train)
predictions = lm.predict(std.transform(X_test))
print("RMSE Linear Regression: ", np.sqrt(mean_squared_error(y_test, predictions)))
list(zip(X_train.columns, lm.coef_))
```

RMSE Linear Regression: 0.06568274378853765

Out[739]:

```
[('GRE Score', 0.02279984761710703),
 ('TOEFL Score', 0.018155915077019393),
 ('University Rating', 0.008078145789796546),
 ('SOP', 0.0009158722244926622),
 ('LOR', 0.012516168309954182),
 ('CGPA', 0.07159838907779793),
 ('Research', 0.009193614982627534)]
```

In [740]:

```
lm = Lasso()
lm.fit(X_train.values, y_train)
predictions = lm.predict(std.transform(X_test))
print("RMSE Lasso Regression: ", np.sqrt(mean_squared_error(y_test, predictions)))
list(zip(X_train.columns, lm.coef_))
```

RMSE Lasso Regression: 0.1448401877933055

Out[740]:

```
[('GRE Score', 0.0),
 ('TOEFL Score', 0.0),
 ('University Rating', 0.0),
 ('SOP', 0.0),
 ('LOR', 0.0),
 ('CGPA', 0.0),
 ('Research', 0.0)]
```

In [741]:

```
lm = Ridge()
lm.fit(X_train.values, y_train)
predictions = lm.predict(std.transform(X_test))
print("RMSE Lasso Regression: ", np.sqrt(mean_squared_error(y_test, predictions)))
list(zip(X_train.columns, lm.coef_))
```

RMSE Lasso Regression: 0.0656776692949077

Out[741]:

```
[('GRE Score', 0.022981409017396162),
 ('TOEFL Score', 0.018297984338214893),
 ('University Rating', 0.008146094998667114),
 ('SOP', 0.0010764691675055765),
 ('LOR', 0.01256553212079061),
 ('CGPA', 0.07094749270353783),
 ('Research', 0.009223428627942506)]
```

## Model Summary using stats model library

In [742]:

```
def build_model(X,y):
    X = sm.add_constant(X) #Adding the constant
    lm = sm.OLS(y.values,X).fit() # fitting the model
    print(lm.summary()) # model summary
    return X

def checkVIF(X):
    vif = pd.DataFrame()
    vif['Features'] = X.columns
    vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
    vif['VIF'] = round(vif['VIF'], 2)
    vif = vif.sort_values(by = "VIF", ascending = False)
    return(vif)
```

## Model 1

In [743]:

```
X_train_new = build_model(X_train, y_train)
```

### OLS Regression Results

```

=====
====
Dep. Variable:          y    R-squared:
0.828
Model:                OLS    Adj. R-squared:
0.825
Method:              Least Squares    F-statistic:          2
70.2
Date:                Thu, 24 Aug 2023    Prob (F-statistic):      1.01e
-145
Time:                22:43:25    Log-Likelihood:          57
1.18
No. Observations:      400    AIC:                      -1
126.
Df Residuals:          392    BIC:                      -1
094.
Df Model:              7
Covariance Type:      nonrobust
=====
=====

```

	coef	std err	t	P> t	[0.025
const	0.7244	0.003	247.175	0.000	0.719
GRE Score	0.0228	0.006	3.676	0.000	0.011
TOEFL Score	0.0182	0.006	3.207	0.001	0.007
University Rating	0.0081	0.005	1.656	0.099	-0.002
SOP	0.0009	0.005	0.186	0.853	-0.009
LOR	0.0125	0.004	2.995	0.003	0.004
CGPA	0.0716	0.006	11.183	0.000	0.059
Research	0.0092	0.004	2.551	0.011	0.002

```

=====
====
Omnibus:              91.100    Durbin-Watson:
1.968
Prob(Omnibus):        0.000    Jarque-Bera (JB):      21
6.020
Skew:                 -1.135    Prob(JB):              1.24
e-47
Kurtosis:             5.794    Cond. No.
5.69
=====
=====

```

#### Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [744]:

```
checkVIF(X_train_new)
```

Out[744]:

	Features	VIF
6	CGPA	4.77
1	GRE Score	4.48
2	TOEFL Score	3.73
4	SOP	2.82
3	University Rating	2.77
5	LOR	2.03
7	Research	1.51
0	const	1.00

### Inference

- p-value of SOP seems to be higher than the significance value of 0.05, hence dropping it as it is insignificant in presence of other variables.

In [745]:

```
X_train_new = X_train_new.drop(["SOP"], axis = 1)
```

## Model 2



In [746]:

```
X_train_new = build_model(X_train_new, y_train)
```

### OLS Regression Results

```
=====
=====
Dep. Variable:          y    R-squared:
0.828
Model:                OLS    Adj. R-squared:
0.826
Method:              Least Squares    F-statistic:          3
16.1
Date:                Thu, 24 Aug 2023    Prob (F-statistic):      5.55e
-147
Time:                22:43:25    Log-Likelihood:          57
1.16
No. Observations:      400    AIC:          -1
128.
Df Residuals:          393    BIC:          -1
100.
Df Model:              6
Covariance Type:      nonrobust
=====
=====
```

	coef	std err	t	P> t	[0.025
const	0.7244	0.003	247.479	0.000	0.719
GRE Score	0.0228	0.006	3.678	0.000	0.011
TOEFL Score	0.0182	0.006	3.224	0.001	0.007
University Rating	0.0084	0.004	1.893	0.059	-0.000
LOR	0.0127	0.004	3.162	0.002	0.005
CGPA	0.0718	0.006	11.427	0.000	0.059
Research	0.0092	0.004	2.562	0.011	0.002

```
=====
=====
Omnibus:              90.434    Durbin-Watson:
1.968
Prob(Omnibus):        0.000    Jarque-Bera (JB):          21
3.539
Skew:                 -1.129    Prob(JB):                  4.27
e-47
Kurtosis:             5.778    Cond. No.
5.26
=====
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [747]:

```
checkVIF(X_train_new)
```

Out[747]:

	Features	VIF
5	CGPA	4.61
1	GRE Score	4.48
2	TOEFL Score	3.72
3	University Rating	2.32
4	LOR	1.89
6	Research	1.51
0	const	1.00

In [748]:

```
X_train_new = X_train_new.drop(["University Rating"], axis = 1)
```

### Inference

- p-value of University Rating seems to be higher than the significance value of 0.05, hence dropping it as it is insignificant in presence of other variables.

## MultiCollinearity Check using VIF

In [749]:

```
checkVIF(X_train_new)
```

Out[749]:

	Features	VIF
1	GRE Score	4.47
4	CGPA	4.38
2	TOEFL Score	3.65
3	LOR	1.71
5	Research	1.50
0	const	1.00

### Inferences

- VIF looks fine and hence, we can go ahead with the predictions

In [750]:

```
lm = sm.OLS(y_train.values,X_train_new).fit()  
y_train_admit = lm.predict(X_train_new)
```


## Residual Analysis of the Model

In [751]:

```
# Plot the histogram of the error terms  
fig = plt.figure()  
sns.distplot((y_train - y_train_admit), bins = 20)  
fig.suptitle('Error Terms', fontsize = 20)    # Plot heading  
plt.show()
```

C:\Users\91798\miniconda3\lib\site-packages\seaborn\distributions.py:2619:  
FutureWarning: `distplot` is a deprecated function and will be removed in  
a future version. Please adapt your code to use either `displot` (a figure  
-level function with similar flexibility) or `histplot` (an axes-level fun  
ction for histograms).

warnings.warn(msg, FutureWarning)



### Inferences

- Error terms seem to be approximately normally distributed, so the assumption on the linear modeling seems to be fulfilled.

## Mean of Residuals

In [752]:

```
residuals = (y_train - y_train_admit)  
np.mean(residuals)
```

Out[752]:

0.008274902905840436

In [753]:

```
X_train_new = X_train_new.drop('const',axis=1)
```

**Mean of Residuals: -0.00277305161296889**



In [754]:




```
# Now let's use our model to make predictions.  
# Creating X_test_new dataframe by dropping variables from X_test  
X_test_new = X_test[X_train_new.columns]  
  
# Adding a constant variable  
X_test_new = sm.add_constant(std.fit_transform(X_test_new))  
# Making predictions  
y_pred = lm.predict(X_test_new)
```

## Linearity of Variables: Residual Plot

In [755]:

```
for i in X_train_new.columns:  
    plt.figure(figsize = (15, 8))  
    sns.residplot(x=X_train_new[i], y=y_train_admit)
```






In [756]:

```
residuals = y_test - y_pred
p = sns.scatterplot(x=y_pred,y=residuals)
plt.xlabel('predicted values')
plt.ylabel('Residuals')
# plt.ylim(-0.4,0.4)
# plt.xlim(0,1)
```

Out[756]:

Text(0, 0.5, 'Residuals')



In [757]:

```
import statsmodels.stats.api as sas
from statsmodels.compat import lzip
name=['F statistics','p-value']
test=sas.het_goldfeldquandt(residuals,X_test)
lzip(name,test)
```

Out[757]:

```
[('F statistics', 1.1495491734096537), ('p-value', 0.32487005690653364)]
```

## Inferences

- Here null hypothesis is - error terms are homoscedastic and since p-values > 0.05, we fail to reject the null hypothesis


## Normality of Residual

In [758]:

```
# Plotting y_test and y_pred to understand the spread.
fig = plt.figure()
plt.scatter(y_test.values, y_pred)
fig.suptitle('y_test vs y_pred', fontsize=20)           # Plot heading
plt.xlabel('y_test', fontsize=18)                     # X-Label
plt.ylabel('y_pred', fontsize=16)
```

Out[758]:

Text(0, 0.5, 'y\_pred')



## Inferences

- y\_test and y\_pred overlaps for the most of the datapoints

In [ ]:


```
fig = plt.figure()
sns.distplot(residuals, bins = 20)
fig.suptitle('Distribution of Residuals', fontsize = 20)
plt.show()
```



In [760]:

```
import matplotlib.pyplot as plt

#create Q-Q plot with 45-degree line added to plot
fig = sm.qqplot(residuals, line = "r")
plt.show()
```



### Inferences

- QQ Plots suggest majority of the data points fit the regression line.

## Model performance evaluation

### Metrics checked - MAE, RMSE, R2, Adj R2

In [761]:

```

from sklearn.metrics import mean_squared_error,r2_score,mean_absolute_error

print("=====" * 10)
print('Mean Absolute Error: ',mean_absolute_error(y_test.values,pred))
print("=====" * 10)
print('Root Mean Square Error: ',np.sqrt(mean_squared_error(y_test.values,pred)))
print("=====" * 10)
r2Score = r2_score(y_test, y_pred)
print('R2 Score: ', r2Score)
print("=====" * 10)
aR2Score = 1 - (1-r2Score/(len(y_test)-X_test_new.shape[1]-1))
print('Adjusted. R2 Score: ', r2Score)
print("=====" * 10)

```

```

=====
Mean Absolute Error:  0.6962010480812433
=====
Root Mean Square Error:  0.7210195623986372
=====
R2 Score:  0.7867445264177775
=====
Adjusted. R2 Score:  0.7867445264177775
=====

```

## Performance test Train & Test Dataset

In [762]:

```

print("=====" * 10)
Trainr2Score = r2_score(y_train, y_train_admit)
print('Train R2 Score: ', Trainr2Score)
print("=====" * 10)
Testr2Score = r2_score(y_test, y_pred)
print('Test R2 Score: ', Testr2Score)

```

```

=====
Train R2 Score:  0.8267706543222355
=====
Test R2 Score:  0.7867445264177775

```

In [763]:

```
print(lm.summary())
```

### OLS Regression Results

```
=====
=====
Dep. Variable:          y    R-squared:
0.827
Model:                OLS    Adj. R-squared:
0.825
Method:             Least Squares    F-statistic:          3
76.1
Date:                Thu, 24 Aug 2023    Prob (F-statistic):      1.61e
-147
Time:                22:43:27    Log-Likelihood:        56
9.34
No. Observations:      400    AIC:                -1
127.
Df Residuals:          394    BIC:                -1
103.
Df Model:              5
Covariance Type:      nonrobust
=====
=====
              coef    std err          t      P>|t|      [0.025
0.975]
-----
const          0.7244      0.003    246.672      0.000      0.719
0.730
GRE Score       0.0233      0.006     3.752      0.000      0.011
0.035
TOEFL Score     0.0197      0.006     3.509      0.001      0.009
0.031
LOR             0.0150      0.004     3.911      0.000      0.007
0.023
CGPA            0.0745      0.006    12.119      0.000      0.062
0.087
Research        0.0096      0.004     2.676      0.008      0.003
0.017
=====
=====
Omnibus:          86.890    Durbin-Watson:
1.980
Prob(Omnibus):    0.000    Jarque-Bera (JB):      20
1.597
Skew:             -1.093    Prob(JB):              1.67
e-44
Kurtosis:         5.705    Cond. No.
4.82
=====
=====
```

#### Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

# Actionable Insights and Recommendations

1. R-squared and Adjusted R-squared (extent of fit) - 0.83 and 0.82 - 85% variance explained.
2. F-stats and Prob(F-stats) (overall model fit) - 387.9 and 1.03e-149(approx. 0.0) - Model fit is significant and explained 82% variance is just not by chance.
3. p-values - p-values for all the coefficients seem to be less than the significance level of 0.05. - meaning that all the predictors are statistically significant.
4. There is lot of chance for the model improvement by tuning the parameters.
5. Currently this models attains accuracy around 80%. This can be improved further by doing some feature engg.
6. As the dataset is strictly provided for the Indian perspective. This model is not generalized, there is scope for the generalization of this model.
7. LogLikelihood is around 570 which indicates model is significantly fit.
8. Performance of training and test data is almost same indicates the model will work significantly on unseen data.
9. While observing the model and according to test assumptions - We can infer errors are homoscedasticity according to p-value
10. While observing the linearity of residual there is no significant pattern found which indicates the residual plots are not correlated
11. While observing the normality of residual - the distribution resembles like bell-shaped and the reg. line fits almost every point

In [ ]: